

Report : Project #2 Continuous Control

5/4/2019

Soyoung Park

Jupyter notebook can be viewed here:

https://nbviewer.jupyter.org/github/parksoy/Soyoung_Udacity_ND_DeepReinforcementLearning/blob/master/p2_continuous-control/Continuous_Control.ipynb

This project is built with the multi-agent [Reacher](#) environment (Unity Machine Learning Agents, open-source Unity plugin that enables games and simulations to serve as environments for training intelligent agents). In this environment, a double-jointed arm can move to target locations. A reward of **+0.1** is provided for each step that the agent's hand is in the goal location. The goal of your agent is to **maintain its position at the target location for as many time steps as possible**. The observation space consists of 33 variables corresponding to position, rotation, velocity, and angular velocities of the arm. Each action is a vector with four numbers, corresponding to torque applicable to two joints. Every entry in the action vector should be a number between -1 and 1.

The version chosen for this project contains 20 identical agents, each with its own copy of the environment. The barrier for solving the multi-agent version of the environment is slightly different, to take into account the presence of many agents. In particular, your agents must get an average score of +30 (over 100 consecutive episodes, and over all agents). Specifically, after **each episode**, we add up the rewards that each agent received (without discounting), to get a score for each agent. This yields 20 (potentially different) scores. We then take the average of these 20 scores. This yields an **average score** for each episode (where the average is over all 20 agents). The environment is considered solved, when **the average (over 100 episodes) of those average scores is at least +30**.

1. Learning Algorithm

- **Learning algorithm**

Deep Deterministic Policy Gradients (DDPG) algorithm ([DDPG paper](#)) was used and its template code from the Actor-Critic Methods lesson was modified to adapt to run Reacher multi-arm version. The various hyperparameters and settings were tweaked.

- **The chosen hyperparameters**

Attempt	max_steps	batch_size	actor_learn_rate	critic_learn_rate	update_every	num_updates	tau	gamma	noise_sigma	noise_factor_decay	layer_sizes
24	1000	1024	0.001	0.001	20	10	0.005	0.99	0.2	1.00E+06	[128,128]

- `max_steps`
: In one episode, how many maximum number of steps to be explored to move onto the next step
- `num_episodes`
: How many times to iterate of training of `max_steps`
- `batch_size`
: How many Replay buffer memory to be sequentially to learn through the neural networks in a batch processing manner
- `buffer_size`
: Replay buffer size that consists of `namedtuple("Experience", field_names=["state", "action", "reward", "next_state", "done"])`
- `actor_learn_rate`
: How fast Actor neural network learns
- `critic_learn_rate`
: How critic Actor neural network learns
- `update_every`
: Update occurs only when certain number of steps is reached to learn more stably.
- `num_updates`
: For every certain step that is defined by `update_every`, only limited number of update is made.
- `tau`
: Control soft update. Only portion of local network is updated, while the majority of update comes from the stable target.
- `gamma`
: Discount factor. Future reward is less valuable than immediate reward.
- `noise_sigma`
: Amount of noise user wants to introduce.
- `layer_sizes`
: Unit size of the first and second hidden layer
- `noise_factor_decay`
: Controls how fast the noise decays over episodes as agents learn the process.
- **Model architectures for neural networks.**

Actor network

```
fc1 = nn.Linear(state_size, fc1_units)
fc2 = nn.Linear(fc1_units, fc2_units)
fc3 = nn.Linear(fc2_units, action_size)
```

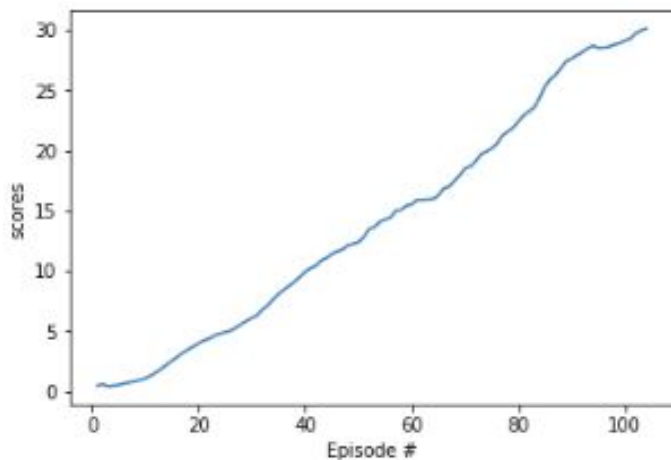
Critic network

```
fcs1 = nn.Linear(state_size, fcs1_units)
fc2 = nn.Linear(fcs1_units+action_size, fc2_units)
fc3 = nn.Linear(fc2_units, 1)
```

where state_size=33, action_size=4, fc1_units=400, fc2_units=300

2. Plot of Rewards: Optimize the hyperparameters

Many attempts were made with different hyperparameters, different tricks to sample/learn to solve the environment as shown in the following table. The final attempt #24 as shown in the following figure demonstrated that 104 episodes were needed to solve the environment (i.e. to receive an average reward 30 points) over all 20 agents.



Atte mpt	max _ste ps	batc h_siz e	actor_le arn_rate	critic_l earn_r ate	upd ate_ ev ery	nu m_ upd ate s	tau	gam ma	nois e_si gma	noise_fact or_decay	layer _size s	score @75 episo de	what's changed
1	300	256	0.0005	0.001	10	10	0.0005	0.99	0.2	0.000001	[128, 128]	0.25	POR
2	300	256	0.0005	0.001	20	10	0.0005	0.99	0.2	0.000001	[128, 128]	0.2	increased buffer size from 3e5 to 3e6 update_every from 10 to 20
3	300	256	0.0005	0.001	20	10	0.0005	0.99	0.2	0.000001	[128, 128]	0.12	torch.nn.utils.clip_grad_norm(self.critic_local.parameters(), 1)
4	1000	256	0.001	0.001	20	10	0.005	0.99	0.2	0.000001	[128, 128]	0.2	to learn faster; tau or soft update of target parameters= 0.0005 to 0.005 actor_learn_rate=0.0005 to 0.001 max_steps=300 to 1000
5	1000	256	0.001	0.001	20	10	0.005	0.99	0.2	0.000001	[128,	0.2	decay noise: actions += self.epsilon *


											128]		self.noise.sample() #decay noise
6	1000	1024	0.001	0.001	20	10	0.005	0.99	0.2	0.000001	[128, 128]	0.2	batch_size=256 to 1024
7	1000	1024	0.001	0.001	10	10	0.005	0.99	0.2	0.000001	[128, 128]	0.15	update_every=20 to 10
8	1000	1024	0.001	0.001	10	10	0.005	0.99	0.2	0.000001	[128, 128]	1.5	for _ in range(num_updates): self.learn(experiences, GAMMA)
9	1000	1024	0.001	0.001	20	10	0.005	0.99	0.2	0.000001	[128, 128]	10	UPDATE EVERY=10 to 20 num_updates=10
10	1000	1024	0.001	0.001	20	10	0.005	0.99	0.2	0.000001	[128, 128]	2.5	avescore_allagents_oneepisode = np.mean(scores) scores_aveofallagents_5episode_deque.append(avescore_allagents_oneepisode) avescore_allagents_5episode = np.mean(scores_aveofallagents_5episode_deque) scores_allagent_allepisode_list.append(avescore_allagents_5episode) avescore_allagents_allepisodes_list.append(avescore_allagents_5episode)
11	1000	512	0.001	0.005	20	10	0.005	0.99	0.2	0.000001	[128, 128]	1.21	max_steps=300 -->max_steps=1000 actor_learn_rate=0.001 critic_learn_rate=0.001-> critic_learn_rate=0.005 batch_size=1024 ->batch_size=512
12	300	128	0.001	0.005	20	20	0.005	0.99	0.2	0.000001	[128, 128]	0.23	max_steps=1000 max_steps=300 update_every=20 num_updates=10->20 batch_size=512->batch_size=128
13	300	512	0.001	0.005	20	10	0.005	0.99	0.2	1.00E+06	[128, 128]	0.32	num_updates=20 to10 batch_size=128 -> batch_size=512
14	300	1024	0.001	0.005	20	10	0.005	0.99	0.2	1.00E+06	[128, 128]	0.2	batch_size=512 -> batch_size=1024 num_updates=20 ->10
15	300	128	0.001	0.005	20	10	0.005	0.99	0.5	1.00E+06	[128, 128]	0.11	NOISE_SIGMA=0.2-> 0.5
16	300	128	0.001	0.005	20	10	0.005	0.99	0.2	1.00E+06	[128, 128]	0.22	noise_sigma=0.5->0.2 turned off batchnorm
17	300	1024	0.001	0.005	20	10	0.005	0.99	0.2	1.00E+06	[128, 128]	0.15	batch_size=128 ->1024
18	300	128	0.001	0.001	20	10	0.005	0.99	0.2	1.00E+06	[128, 128]	0.13	critic_learn_rate=0.005->0.001
19	300	128	0.001	0.001	20	10	0.001	0.99	0.2	1.00E+06	[128, 128]	0.15	tau=0.005 ->tau=0.001
20	300	128	0.001	0.005	20	10	0.005	0.99	0.2	1.00E+06	[400, 300]	0.16	layer_sizes=[128,128] ->[400,300]
21	300	128	0.001	0.005	20	10	0.005	0.99	0.2	1.00E+06	[128, 128]	0.18	removed for _ in range(num_updates):
22	300	128	0.001	0.005	20	10	0.005	0.99	0.2	1.00E+06	[128, 128]	0.2	fixed scores to score
23	300	128	0.001	0.005	20	10	0.005	0.99	0.2	1.00E+06	[128, 128]	0.16	fixed scores to score
24	1000	1024	0.001	0.001	20	10	0.005	0.99	0.2	1.00E+06	[128, 128]	18	Final: repeat attempt # 9 hyperparameters

The key breakthrough occurred at Attempt #8, #9. In agent_reacher.py, two major points were helpful to start learning.


```
def step(self, states, actions, rewards, next_states, dones, t):
    for state, action, reward, next_state, done in zip(states, actions, rewards, next_states, dones):
        self.memory.add(state, action, reward, next_state, done)

    if len(self.memory) > BATCH_SIZE and t % UPDATE_EVERY == 0:
        experiences = self.memory.sample()
        for _ in range(num_updates):
            experiences = self.memory.sample()
            self.learn(experiences, GAMMA)
```

In Attempt #8, `for _ in range(num_updates):` was added so learning takes only less than `num_updates` times (hyperparameter, set as 10) every `UPDATE_EVERY` (set as 10) step.



In Attempt #9, `UPDATE_EVERY` was updated from 10 to 20 to learn more stably. Agents step and memory is filled and memory is sampled every `UPDATE_EVERY` (increased from 10 to 20) but learning occurs only `num_updates`(10) times.



3. Ideas for Future Work

For improving the agent's performance, the various deep RL algorithms on continuous control tasks, such as, REINFORCE, TNPG, RWR, REPS, TRPO, CEM, CMA-ES and DDPG, may be pursued based on suggestion in the review [paper](#).